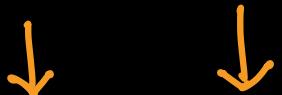




Integration Delivery / Deployment



CI/CD



What is CI/CD Pipelines

→ Sequence of stages → SCM, build, test, config, deploy



- A CI/CD pipeline automates the steps required from writing code - delivering the application - deploying it to production
- It ensures:
 - Faster development cycles
 - Higher code quality
 - Fewer manual errors
 - Repeatable and reliable deployments
- Benefits of CI/CD pipeline
 - Business Benefits
 - Faster time to market
 - Less downtime
 - Lower cost of development
 - Increased reliability → automation
 - Developer Benefits
 - Stable codebase
 - Automated workflows
 - Faster feedback loop
 - Easier collaboration

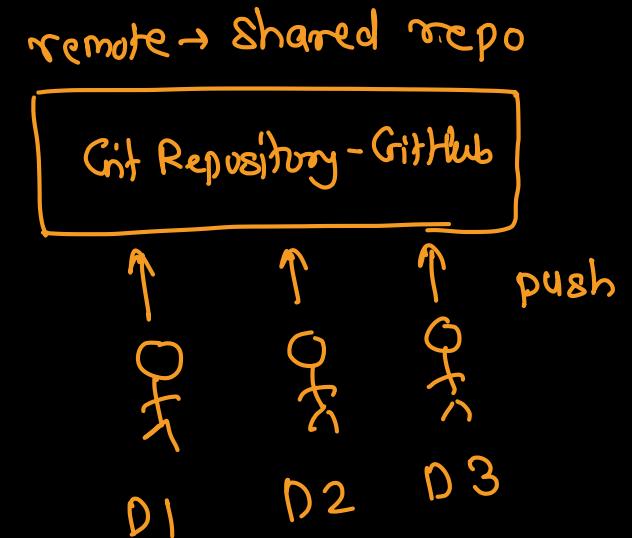
Continuous Integration



- This is the practice of automatically integrating code changes from multiple developers into a shared repository frequently—often several times a day

- What happens:

- Developers push code to a version control system (like Git)
- Automated tests run to check if the new code breaks anything
- The code is automatically built/compiled
- Teams get immediate feedback if something is wrong





The Problem CI Solves

▪ Before CI

- Developers worked on separate branches for weeks or months
- When merging, massive conflicts occurred ("integration hell")
- Bugs were discovered late, making them expensive to fix
- No one knew if the combined code actually worked until the end
- Building and testing were manual, slow, and error-prone

▪ With CI

- Small, frequent integrations reduce conflict complexity
- Problems are detected within minutes, not weeks
- The codebase is always in a working state → main branch
- Team collaboration improves with better visibility → Dashboard

Core Principles of CI



Single Source Repository (Remote)

- All code lives in one version control system (Git, SVN, etc.). Everyone commits to the same repository, typically using a branching strategy like GitFlow or trunk-based development.

Frequent Commits

- Developers commit code at least once per day, ideally multiple times. Smaller, incremental changes are easier to integrate and debug than large batches.

Automated Build

- Every commit triggers an automated build process that compiles the code and creates executable artifacts. This ensures the code can actually be built successfully.

→ deployable package

Self-Testing Build

- The build includes running automated tests. If tests fail, the build is considered broken. Common test types include:
- Unit tests - Test individual functions/components
- Integration tests - Test how components work together
- Static code analysis - Check code quality and standards
- Security scans - Identify vulnerabilities

Fast Builds

- Builds should complete quickly (ideally under 10 minutes) so developers get rapid feedback. Slow builds discourage frequent commits.

Test in Production-Like Environment → pre-production / UAT

- CI environments should mirror production as closely as possible to catch environment-specific issues early.

Easy Access to Latest Build

- Everyone on the team can access the latest executable and test results. Transparency is crucial.

Visible Build Status → Dashboard

- The current build status is highly visible to the team through dashboards, notifications, or physical indicators (like lava lamps or traffic lights in some offices).

Fix Broken Builds Immediately

- When a build breaks, fixing it becomes the top priority. The team shouldn't commit new changes until the build is green again.

The CI Workflow



■ Step 1: Developer Works Locally

- Developer pulls the latest code from the main branch
- Makes changes and writes tests → TDD
- Runs tests locally to ensure everything works

■ Step 2: Commit and Push

- Developer commits changes with a descriptive message
- Pushes code to the central repository

■ Step 3: CI Server Detects Change

- The CI server (Jenkins, GitLab CI, etc.) monitors the repository
- Detects the new commit via webhooks or polling

■ Step 4: Checkout Code

- CI server pulls the latest code from the repository
- Creates a clean build environment

■ Step 5: Build Process

- Compiles the source code
- Resolves dependencies
- Creates build artifacts (executables, containers, packages)

■ Step 6: Run Tests

- Executes automated test suites
- Generates test reports and code coverage metrics
- May run multiple test stages in parallel for speed

■ Step 7: Code Quality Checks

- Runs linters to check coding standards
- Performs static code analysis → SonarQube
- Checks for security vulnerabilities ↗ code coverage
- Verifies code coverage meets thresholds

■ Step 8: Generate Artifacts

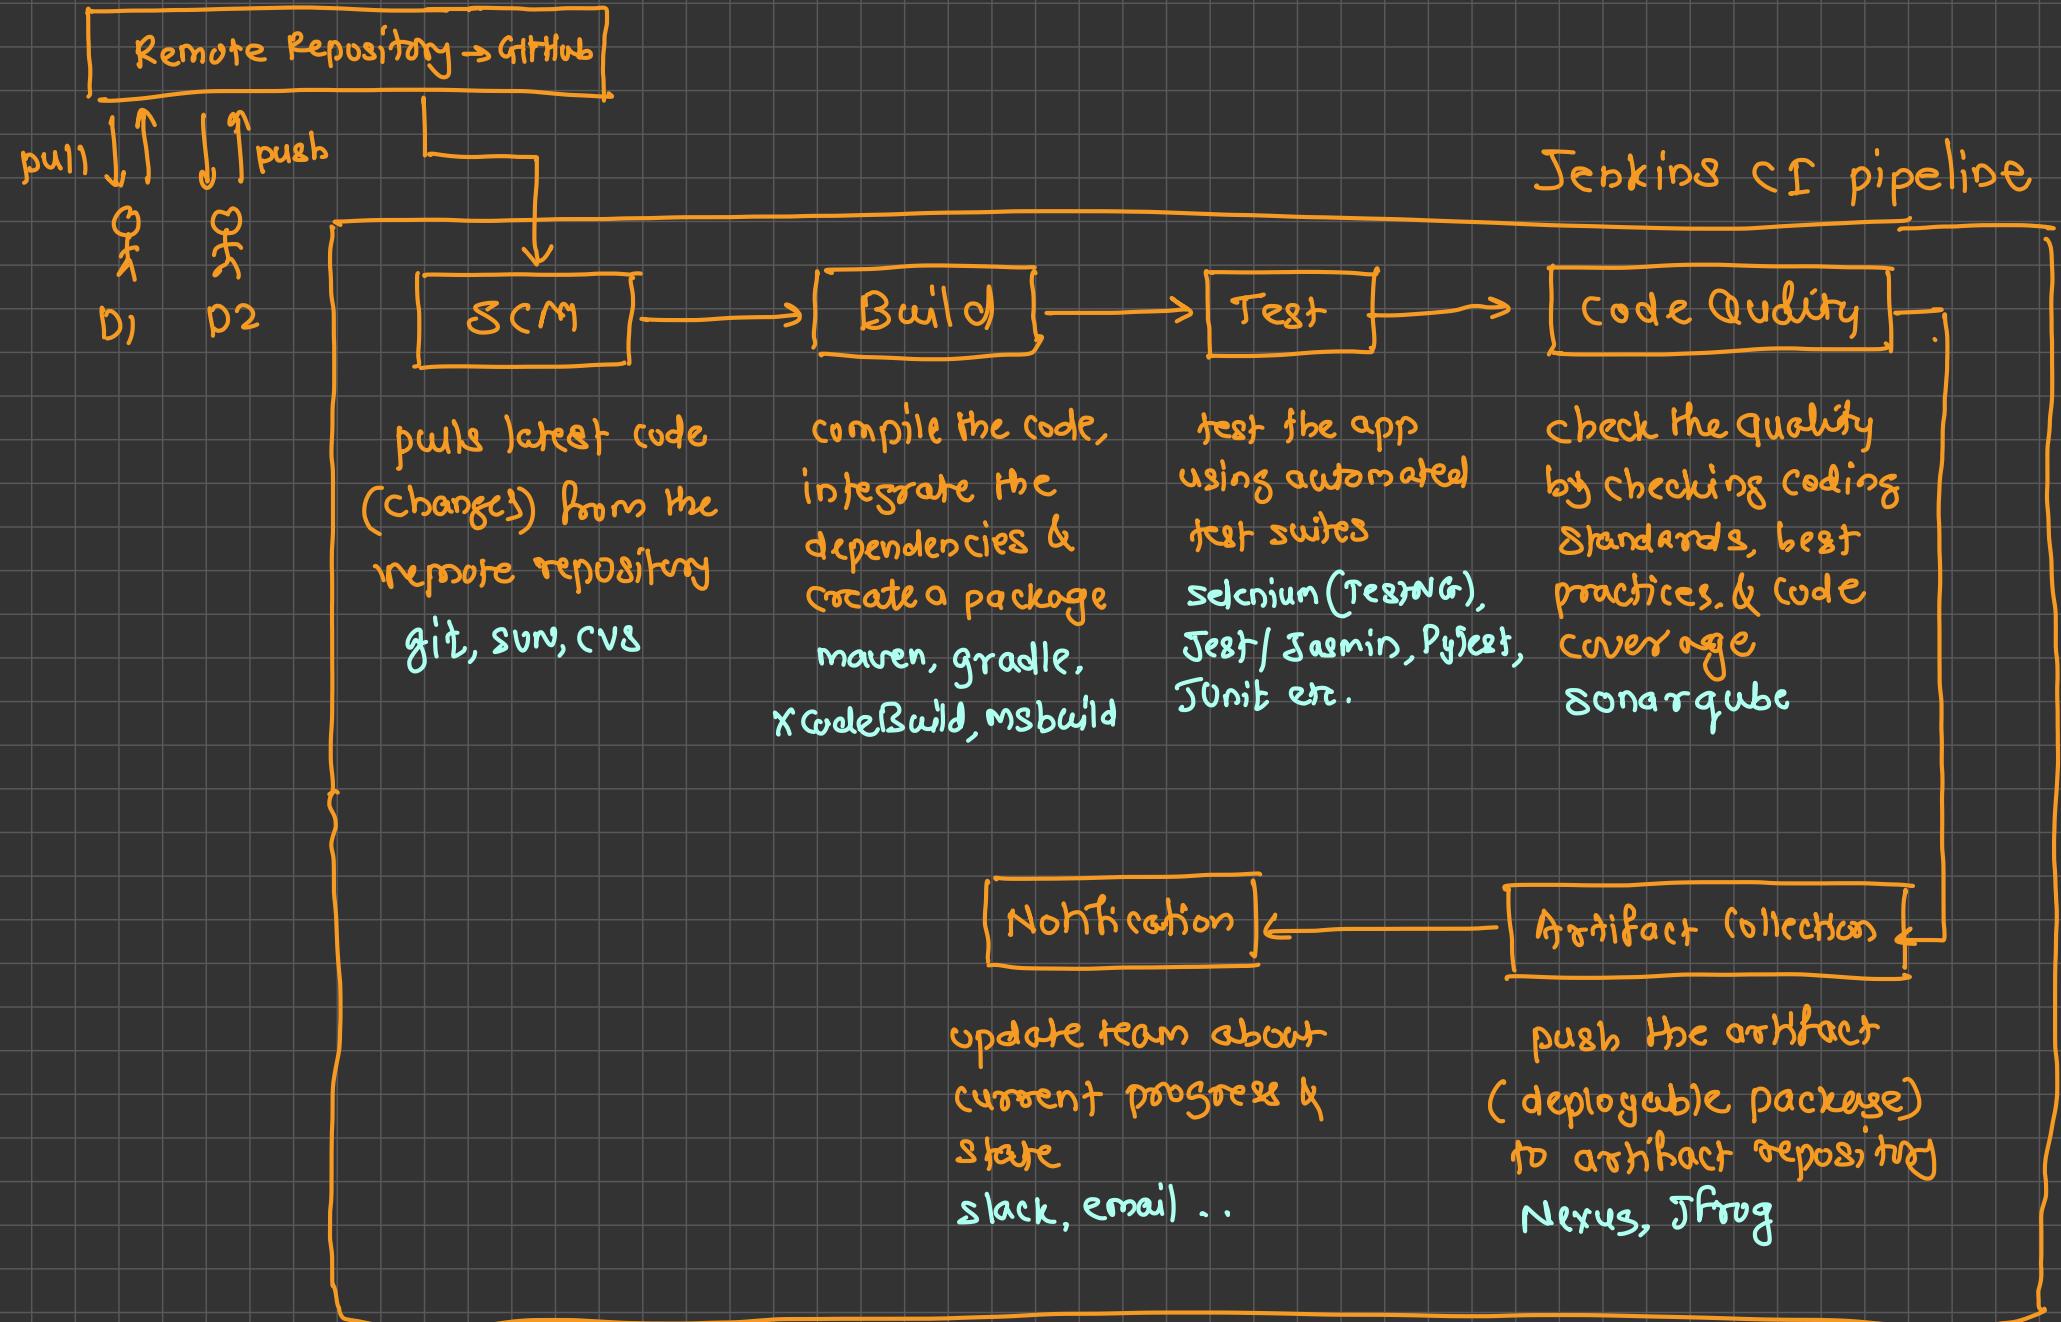
- If all checks pass, creates deployable artifacts
- Tags the build with version numbers
- Stores artifacts in a repository (Artifactory, Nexus, container registry)

■ Step 9: Notification

- Sends results to the team (email, Slack, dashboard)
- Updates build status badges
- Green = success, Red = failure

■ Step 10: Feedback Loop

- If successful: Developer continues with next task
- If failed: Developer investigates logs, fixes issues, and recommits



Key CI Components



▪ Version Control System (VCS)

- Git (GitHub, GitLab, Bitbucket)
- Subversion → SVN
- Mercurial

▪ CI Server/Platform

- ✓ Jenkins - Most popular, highly customizable, open-source
- GitLab CI/CD - Integrated with GitLab, uses YAML configs
- GitHub Actions - Native to GitHub, marketplace of actions
- CircleCI - Cloud-based, fast Docker support
- Travis CI - Popular for open-source projects
- TeamCity - By JetBrains, powerful and user-friendly
- Azure DevOps - Microsoft's comprehensive platform
- Bamboo - By Atlassian, integrates with Jira

▪ Build Tools

- Maven, Gradle (Java)
- npm, Yarn, Webpack (JavaScript)
- MSBuild (.NET)
- Make, CMake (C/C++)
- pip, Poetry (Python)

▪ Build Tools

- Maven, Gradle (Java)
- npm, Yarn, Webpack (JavaScript)
- MSBuild (.NET)
- Make, CMake (C/C++)
- pip, Poetry (Python)

▪ Testing Frameworks

- JUnit, TestNG (Java)
- Jest, Mocha, Cypress (JavaScript)
- pytest, unittest (Python)
- NUnit, xUnit (.NET)

▪ Code Quality Tools

- SonarQube - Comprehensive code analysis
- ESLint, Pylint - Linters for specific languages
- CodeClimate, Codacy - Cloud-based analysis
- OWASP Dependency Check - Security scanning



Benefits of CI

- Reduced Integration Risk - Small, frequent merges prevent "integration hell"
- Higher Code Quality - Automated testing catches bugs early
- Faster Feedback - Developers know within minutes if their code works
- Improved Collaboration - Visibility into what everyone is working on
- Reduced Manual Testing - Automation frees up time for exploratory testing
- Documentation - Build logs provide a history of changes and issues
- Confidence - Teams can refactor and improve code without fear
- Faster Time to Market - Streamlined process accelerates delivery



Continuous Delivery and Continuous Deployment

manual deployment

automated Deployment

- CD refers to two related but distinct practices that extend Continuous Integration:
- Continuous Delivery** - Code changes are automatically built, tested, and prepared for release to production, but deployment requires manual approval. → manual deployment
- Continuous Deployment** - Every change that passes all automated tests is automatically released to production without human intervention. → automated deployment
- Both ensure that code is always in a deployable state, but they differ in the final step to production.

The Problem CD Solves



▪ Before CD

- Releases happened infrequently (monthly, quarterly, or even yearly)
- Manual deployment processes were error-prone and stressful
- Large releases contained many changes, making issues hard to trace
- Deployment knowledge was siloed with specific team members
- Testing in production-like environments happened too late
- Rollbacks were complex and risky
- Long lead times from code completion to customer value

▪ With CD

- Releases can happen multiple times per day
- Deployments are automated, consistent, and repeatable
- Small, incremental changes reduce risk
- Anyone on the team can deploy with a button click
- Issues are detected early through extensive automated testing
- Rollbacks are quick and automated
- Features reach customers rapidly

Continuous Delivery vs Continuous Deployment



▪ Continuous Delivery

- Manual gate before production deployment
- Humans decide when to release
- Business decides the release schedule
- Suitable when releases need coordination (marketing, training, etc.) ↗ internal
- Common in enterprise, regulated industries, or when deploying to client-managed environments
- Example Flow: Code → Build → Test → Stage → [Manual Approval] → Production

▪ Continuous Deployment

- Fully automated pipeline to production
- Every successful build goes live automatically
- No human intervention in the deployment process
- Requires mature testing and monitoring
- Common in SaaS products, web applications, and companies like Netflix, Amazon, Etsy
- Example Flow: Code → Build → Test → Stage → Production (all automated)

Core Principles of CD



- **Build Quality In**
 - Quality isn't checked at the end; it's built into every step. Automated testing at multiple levels ensures confidence in releases.
- **Work in Small Batches**
 - Deploy small, incremental changes frequently rather than large batches. Smaller changes are easier to test, deploy, and troubleshoot.
- **Automate Everything**
 - From building to testing to deployment, minimize manual steps. Automation ensures consistency and reduces human error.
- **Relentless Pursuit of Continuous Improvement**
 - Continuously optimize the pipeline, reduce cycle time, and improve feedback loops.
- **Everyone is Responsible**
 - The entire team owns the deployment pipeline, not just operations. Developers are accountable for production readiness.
- **Done Means Released**
 - Work isn't complete when code is written—it's complete when it's delivering value in production.



The CD Pipeline Stages

▪ **Stage 1: Source Control**

- Everything starts with code committed to version control. This triggers the entire pipeline.

▪ **Stage 2: Build Automation**

- Compile source code
- Resolve and download dependencies
- Create build artifacts (JAR files, Docker images, executables)
- Version and tag artifacts
- Store in artifact repository
- Tools: Maven, Gradle, npm, Docker, Artifactory, Nexus

▪ **Stage 3: Automated Testing**

- This is the most critical part of CD. Multiple layers of testing provide confidence:
- Unit Tests
- Integration Tests
- Contract Tests
- End-to-End Tests
- Performance Tests
- Security Tests

The CD Pipeline Stages



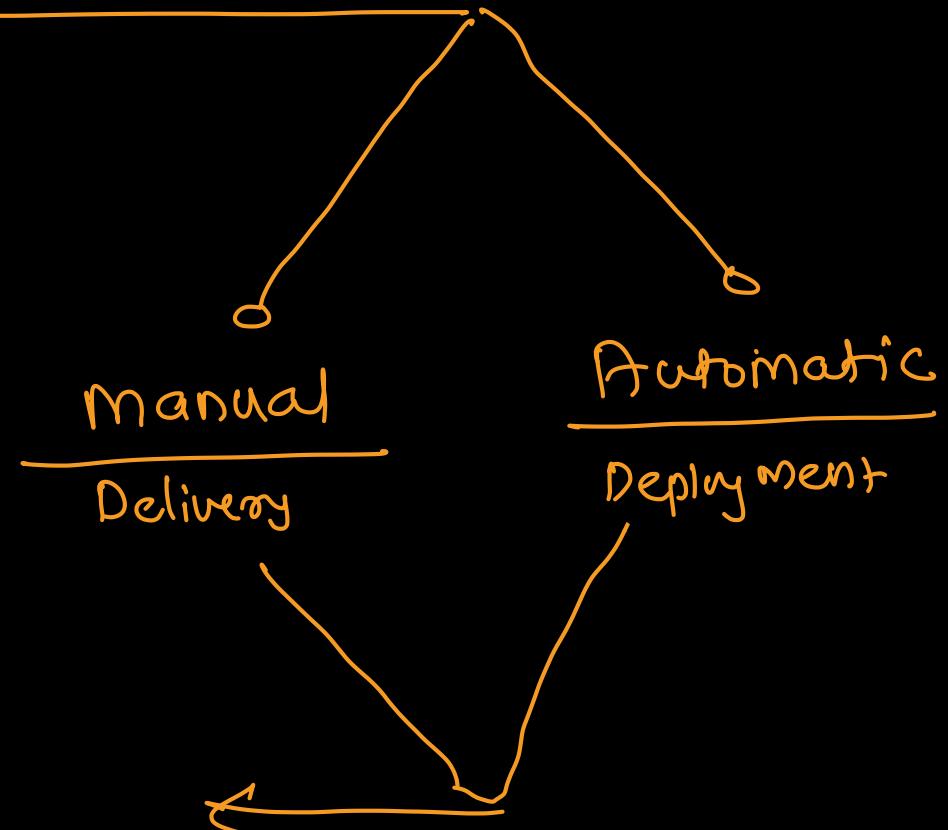
■ Stage 4: Staging/Pre-Production Deployment

- Code is deployed to an environment that mirrors production:
 - Same infrastructure configuration
 - Similar data volumes (often anonymized production data)
 - Same software versions
 - Equivalent network topology
- Activities:
 - Deploy using the same process as production
 - Run full test suites
 - Perform exploratory testing
 - Validate configuration
 - Check monitoring and logging

■ Stage 5: Production Deployment

- The final deployment to live environment where users access the application
- Deployment Strategies:
 - Blue-Green Deployment
 - Canary Deployment
 - Rolling Deployment

→ mirrors of production





Benefits of CD

- **Faster Time to Market** - Features reach customers quickly
- **Reduced Risk** - Small, frequent deployments are less risky
- **Higher Quality** - Extensive automated testing catches issues
- **Lower Stress** - Deployments become routine, not events
- **Better Feedback** - Quick feedback from real users
- **Improved Developer Productivity** - Less time spent on manual processes
- **Competitive Advantage** - Respond to market changes rapidly
- **Lower Costs** - Automation reduces manual effort
- **Better Security** - Patches and fixes deployed quickly
- **Improved Reliability** - Consistent, tested deployment processes



Key CD Technologies and Tools

Infrastructure as Code (IaC)

- Define infrastructure in version-controlled files
- ✓ Terraform - Multi-cloud infrastructure provisioning
- CloudFormation - AWS-specific IaC
- ✓ Ansible - Configuration management and provisioning
- Chef, Puppet - Configuration management
- Pulumi - IaC using programming languages

Container Orchestration

- Manage containerized applications
- ✓ Kubernetes - Industry-standard orchestration
- ✓ Docker Swarm - Simpler alternative to K8s
- Amazon ECS/EKS - AWS container services
- Azure AKS - Azure Kubernetes Service
- Google GKE - Google Kubernetes Engine

Deployment Tools

- Automate the deployment process
- ✓ Helm - Kubernetes package manager
- ✓ ArgoCD - GitOps continuous delivery for K8s
- Flux - GitOps toolkit for K8s
- Spinnaker - Multi-cloud continuous delivery
- Octopus Deploy - Deployment automation
- AWS CodeDeploy - AWS deployment service

Configuration Management

- Manage application configuration
- Consul - Service mesh and configuration
- etcd - Distributed key-value store
- Spring Cloud Config - Centralized configuration for Spring
- AWS Systems Manager - Parameter Store and secrets
- Azure Key Vault - Secrets management

Key CD Technologies and Tools



■ Artifact Repositories

- Store build artifacts
- Docker Hub - Container images (public)
- Amazon ECR - AWS container registry
- Google Container Registry - GCP container registry
- Artifactory - Universal artifact repository
- Nexus - Repository manager
- GitHub Packages - Integrated with GitHub

■ Monitoring and Observability

- Track application health:
- Prometheus + Grafana - Metrics and visualization
- Datadog - Full-stack monitoring
- New Relic - APM and observability
- Dynatrace - AI-powered monitoring
- ELK Stack (Elasticsearch, Logstash, Kibana) - Log management
- Splunk - Log analysis and SIEM
- Jaeger, Zipkin - Distributed tracing
- Sentry - Error tracking



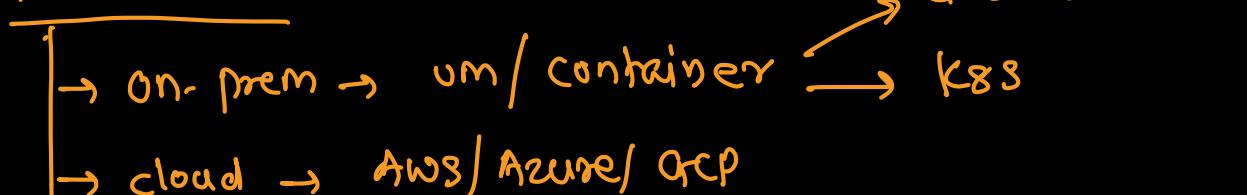
Jenkins



What is Jenkins ? → CI/CD tool free

- Jenkins is an open-source automation server that helps software development teams automate the process of building, testing, and deploying their applications
- Think of Jenkins as a reliable assistant that can perform repetitive tasks automatically, saving developers time and reducing human errors
- It was originally called Hudson and was created in 2004 by Kohsuke Kawaguchi while he worked at Sun Microsystems
- In 2011, it was renamed to Jenkins after a dispute with Oracle, and it has grown to become the most popular automation server in the world
- Jenkins is free to use and has a large community of contributors who continuously improve it and create extensions for it

installation



↳ plugins

Why Jenkins



▪ Strengths

- Free and Open Source - No licensing costs
- Highly Extensible - Over 1,800 plugins available
- Platform Agnostic - Runs on Windows, Linux, macOS
- Technology Agnostic - Works with any programming language or tool
- Large Community - Extensive documentation and community support
- Self-Hosted - Full control over your CI/CD infrastructure
- Mature and Battle-Tested - Used by thousands of organizations
- Flexible - Can be configured for simple to extremely complex workflows

▪ Weaknesses

→ Scripted pipeline → Pipeline as a Code → Groovy

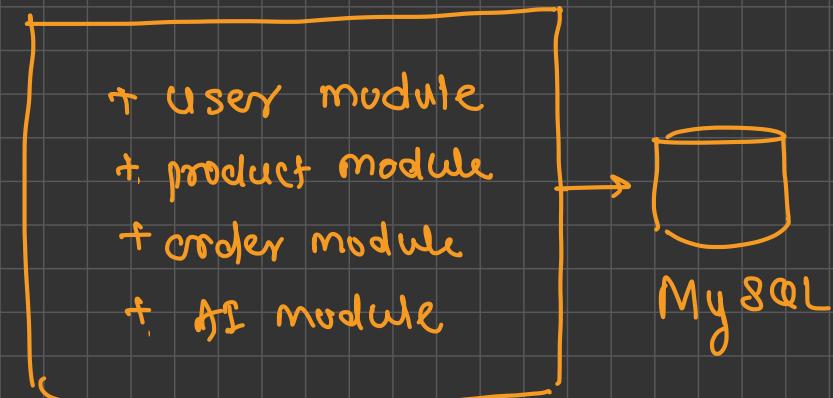
- Steep Learning Curve - Complex to set up and configure initially
- Maintenance Overhead - Requires dedicated resources for upkeep
- UI/UX - Interface feels dated compared to modern alternatives
- Plugin Management - Plugin compatibility issues can arise
- Scaling Challenges - Requires careful architecture for large deployments
- Security - Requires vigilant security configuration and updates

monolithic architecture

- single app has all the features
 - use one language, one technology, one repository & one database
 - less learning curve
 - deployment is easy
 - faster than microservices
 - less skills are required to maintain
- No fault isolation
 - No flexibility for language & Database

Repository

- e-commerce
- user module
- product module
- order module
- AI module



language → Java
framework → Spring

micro-services architecture

→ instead of one big app, MS is used to implement single functionality

→ collection of multiple smaller applications called as MS

→ flexibility on language & DB

→ fault isolation

→ easy scaling

→ complex architecture → slower & costly
→ overhead
→ more skills are required
→ more steep learning curve

Kafka / RabbitMQ message queue

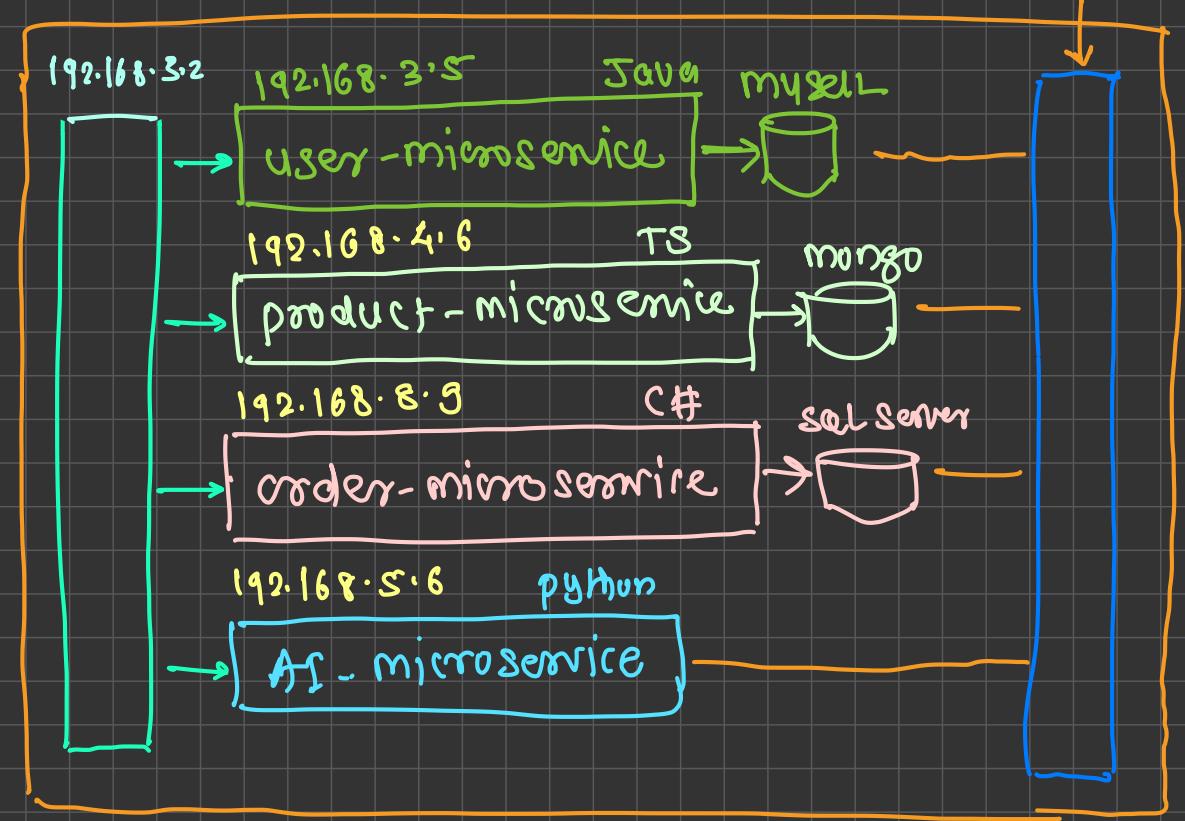
e-commerce

→ user module → Java / MySQL

→ product module → TS / Mongo

→ order module → C# / MS SQL Server

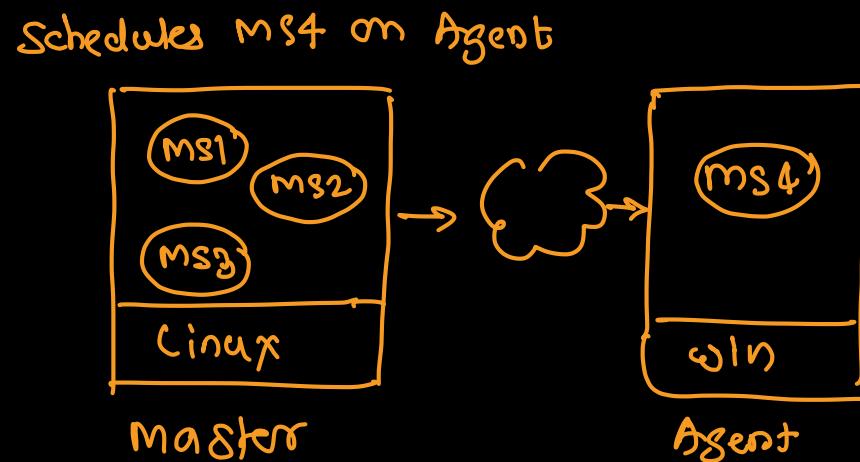
→ AI module → Python



Architecture



- Jenkins operates using a master-agent architecture, which means it has one central controller (the master) and can have multiple workers (agents)
- The master server is the brain of Jenkins - it schedules builds, monitors agents, records results, and provides the web interface you interact with
- Agents are separate machines that actually execute the work - they compile code, run tests, and perform deployments as instructed by the master
- This architecture allows Jenkins to distribute work across multiple machines, enabling it to handle many projects simultaneously
- For small projects, you might run everything on the master, but for larger organizations, you would have dozens or hundreds of agents





Node Types

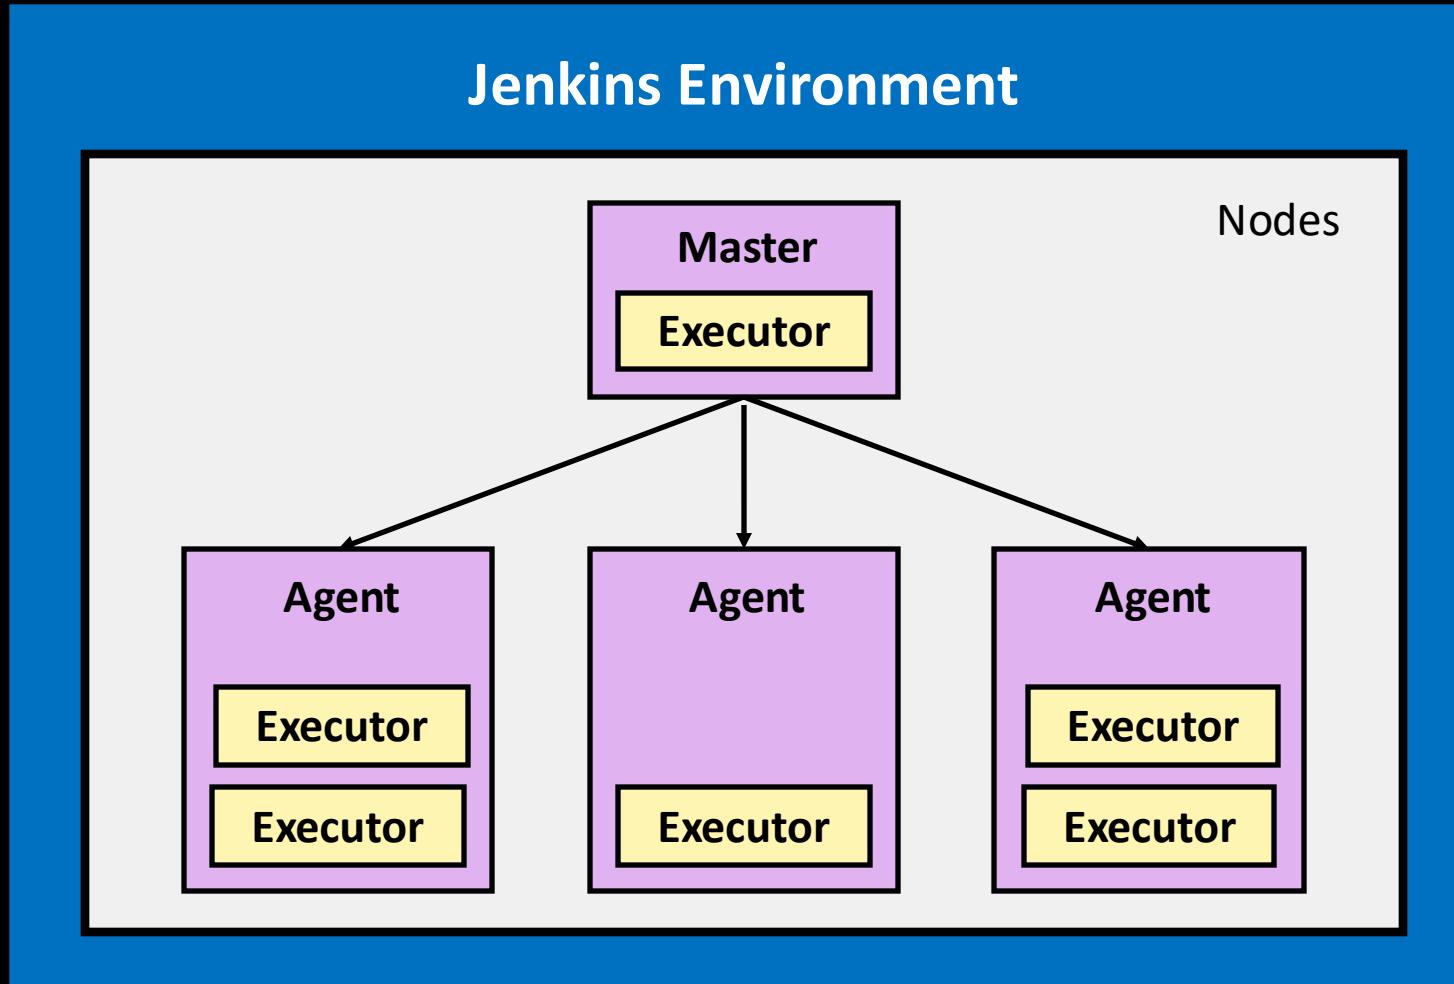
■ Master Server

- The master server is where Jenkins itself runs and where you configure all your automation tasks through a web browser
- It stores all the configuration data, including information about what tasks to run, when to run them, and where to run them
- The master continuously monitors your code repositories (like GitHub or GitLab) for changes and triggers builds when new code is committed
- It also maintains the build history, showing you which builds succeeded, which failed, and providing access to logs and test results
- Think of the master as the conductor of an orchestra, coordinating all the different pieces but not necessarily playing every instrument itself

■ Agent Node

- Agents (also called slaves or nodes) are separate computers or virtual machines that connect to the Jenkins master to perform actual work
- When the master decides a build needs to run, it sends instructions to an available agent, which then executes those instructions
- Agents can run on different operating systems (Windows, Linux, Mac) allowing you to test your application on multiple platforms simultaneously
- Each agent can have labels describing its capabilities, like "linux," "docker," or "high-memory," helping Jenkins choose the right agent for each job
- Using multiple agents allows your organization to run many builds in parallel, dramatically reducing wait times when multiple developers are committing code

Jenkins Environment





- **Freestyle jobs** are the most basic type - they provide a simple web form where you configure each step of your automation process
- **Pipeline jobs** use code (written in a language called Groovy) to define your automation, making it easier to manage complex workflows and version control your automation scripts
- **Multi-configuration jobs** allow you to run the same job with different parameters, like testing your application on Windows, Linux, and Mac simultaneously
- **Folder jobs** help organize your jobs into hierarchies, similar to how you organize files on your computer into folders
- **Multi-branch pipeline jobs** automatically create a separate pipeline for each branch in your source code repository, perfect for teams working on multiple features simultaneously

Jenkins Builds



- A build is a single execution of a Jenkins job - every time a job runs, it creates a new build with an incrementing number
- E.g., if you have a job called "Build My App," its first execution is Build 1, the second execution is Build 2, and so on
- Each build has a status: Success (everything worked), Failure (something went wrong), Unstable (it ran but tests failed), or Aborted (someone stopped it)
- Jenkins keeps a history of all builds, allowing you to see trends over time, like whether your project is getting more stable or less stable
- You can view logs for each build to see exactly what happened during execution, which is invaluable for troubleshooting when something goes wrong

Jenkins Workspace



- A workspace is a directory on an agent's hard drive where Jenkins stores the files needed for a particular job
- When a build starts, Jenkins typically checks out your source code from a repository into the workspace, giving the build access to your files
- The workspace persists between builds of the same job on the same agent, which can speed up subsequent builds by avoiding redundant downloads
- However, this also means that leftover files from previous builds might interfere with new builds, so Jenkins provides options to clean the workspace before each build
- Each job usually has its own workspace directory to prevent different jobs from interfering with each other's files

Jenkins Plugins



- Plugins are extensions that add new functionality to Jenkins, similar to how apps add features to your smartphone
- Jenkins ships with basic functionality, but plugins allow you to integrate with hundreds of different tools and services like GitHub, Docker, Kubernetes, and Slack
- For example, the Git plugin lets Jenkins check out code from Git repositories, while the Email plugin lets Jenkins send notification emails
- There are over 1,800 plugins available in the Jenkins plugin repository, covering everything from source control to deployment to notifications
- Plugins are maintained by the Jenkins community, and you can install them through Jenkins' web interface with just a few clicks

Installing Jenkins - System Requirements



- Before installing Jenkins, your computer or server needs Java installed because Jenkins is a Java application
- Specifically, Jenkins requires Java 11 or Java 17 (these are versions of the Java runtime environment)
- You will need at least 256 MB of RAM, but 1 GB or more is recommended for any real-world usage
- About 10 GB of hard drive space is recommended - 1 GB for the Jenkins application itself and the rest for storing build data and workspaces
- Jenkins can run on Windows, Linux, or Mac operating systems, and there are installation packages available for all major platforms
- **Ubuntu/Debian**
 - `wget -q -O - https://pkg.jenkins.io/debian-stable/jenkins.io.key | sudo apt-key add -`
 - `sudo sh -c 'echo deb https://pkg.jenkins.io/debian-stable binary/ > /etc/apt/sources.list.d/jenkins.list'`
 - `sudo apt update`
 - `sudo apt install jenkins`
- **Start Jenkins**
 - `sudo systemctl start jenkins`
 - `sudo systemctl enable jenkins`