



Kubernetes

Container orchestration Engine



What is Kubernetes ?

→ Google

- Portable, extensible, open-source platform for managing containerized workloads and services
- Facilitates both declarative configuration and automation
- It has a large, rapidly growing ecosystem
- Kubernetes services, support, and tools are widely available
- The name Kubernetes originates from Greek, meaning helmsman or pilot
- Google open-sourced the Kubernetes project in 2014

Ecosystem → no q tools supporting / extending K8S

- observability → Prometheus
- monitoring → Grafana, Kibana
- Logging → Elasticsearch, Logstash
- Service mesh → Istio

declarative config → using YAML syntax

→ simple text file

→ can be tracked by SCM tool

automation → using commands

K8S cluster configuration

→ on-prem

↳ private infrastructure

→ cloud

↳ self managed

→ cloud managed service

→ AWS → EKS → Elastic K8S Service

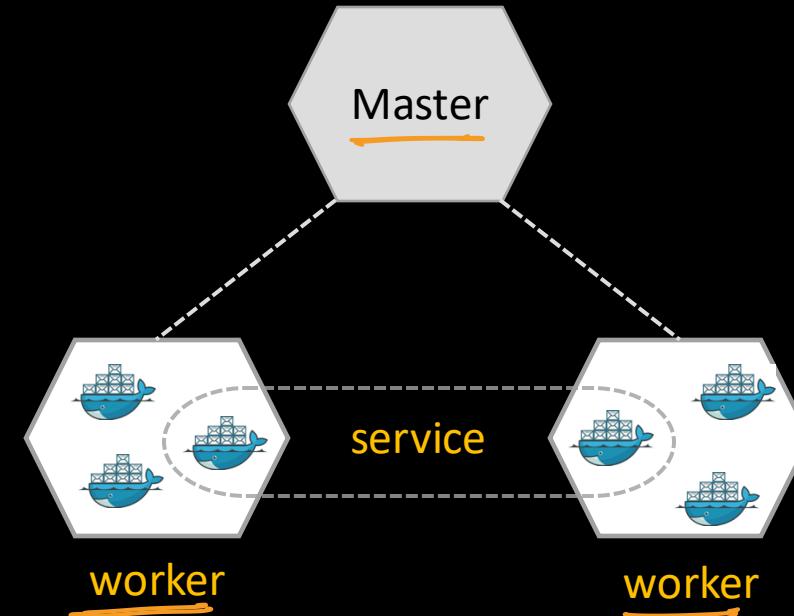
→ Azure → AKS → Azure K8S Service

→ GCP → GKE → GCP K8S Engine



Kubernetes Cluster

- When you deploy Kubernetes, you get a cluster.
- A cluster is a set of machines (nodes), that run containerized applications managed by Kubernetes
- A cluster has at least one worker node and at least one master node
- The worker node(s) host the pods that are the components of the application
- The master node(s) manages the worker nodes and the pods in the cluster
- Multiple master nodes are used to provide a cluster with failover and high availability → prod env



K8S cluster

→ single node cluster

- minikube → fake cluster
- Used Only for learning K8S

→ multi-node cluster

→ Single master cluster

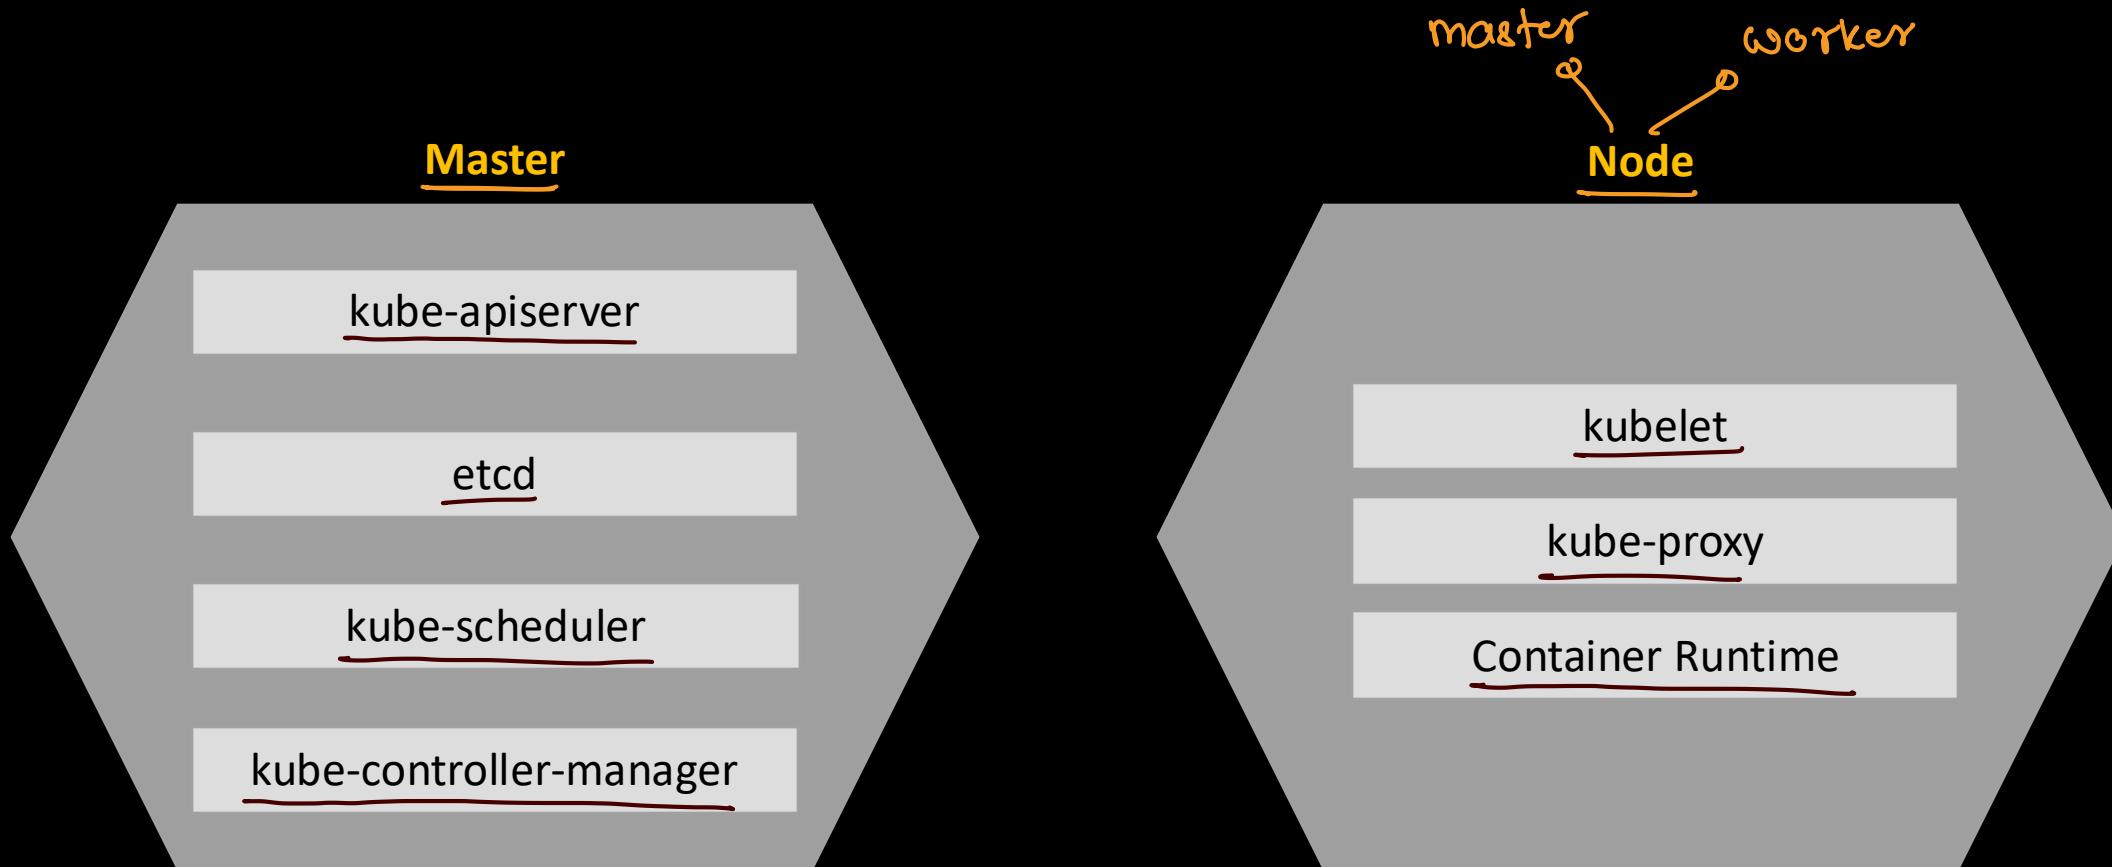
- only one master and one or more workers
- used in dev and staging/ testing env.

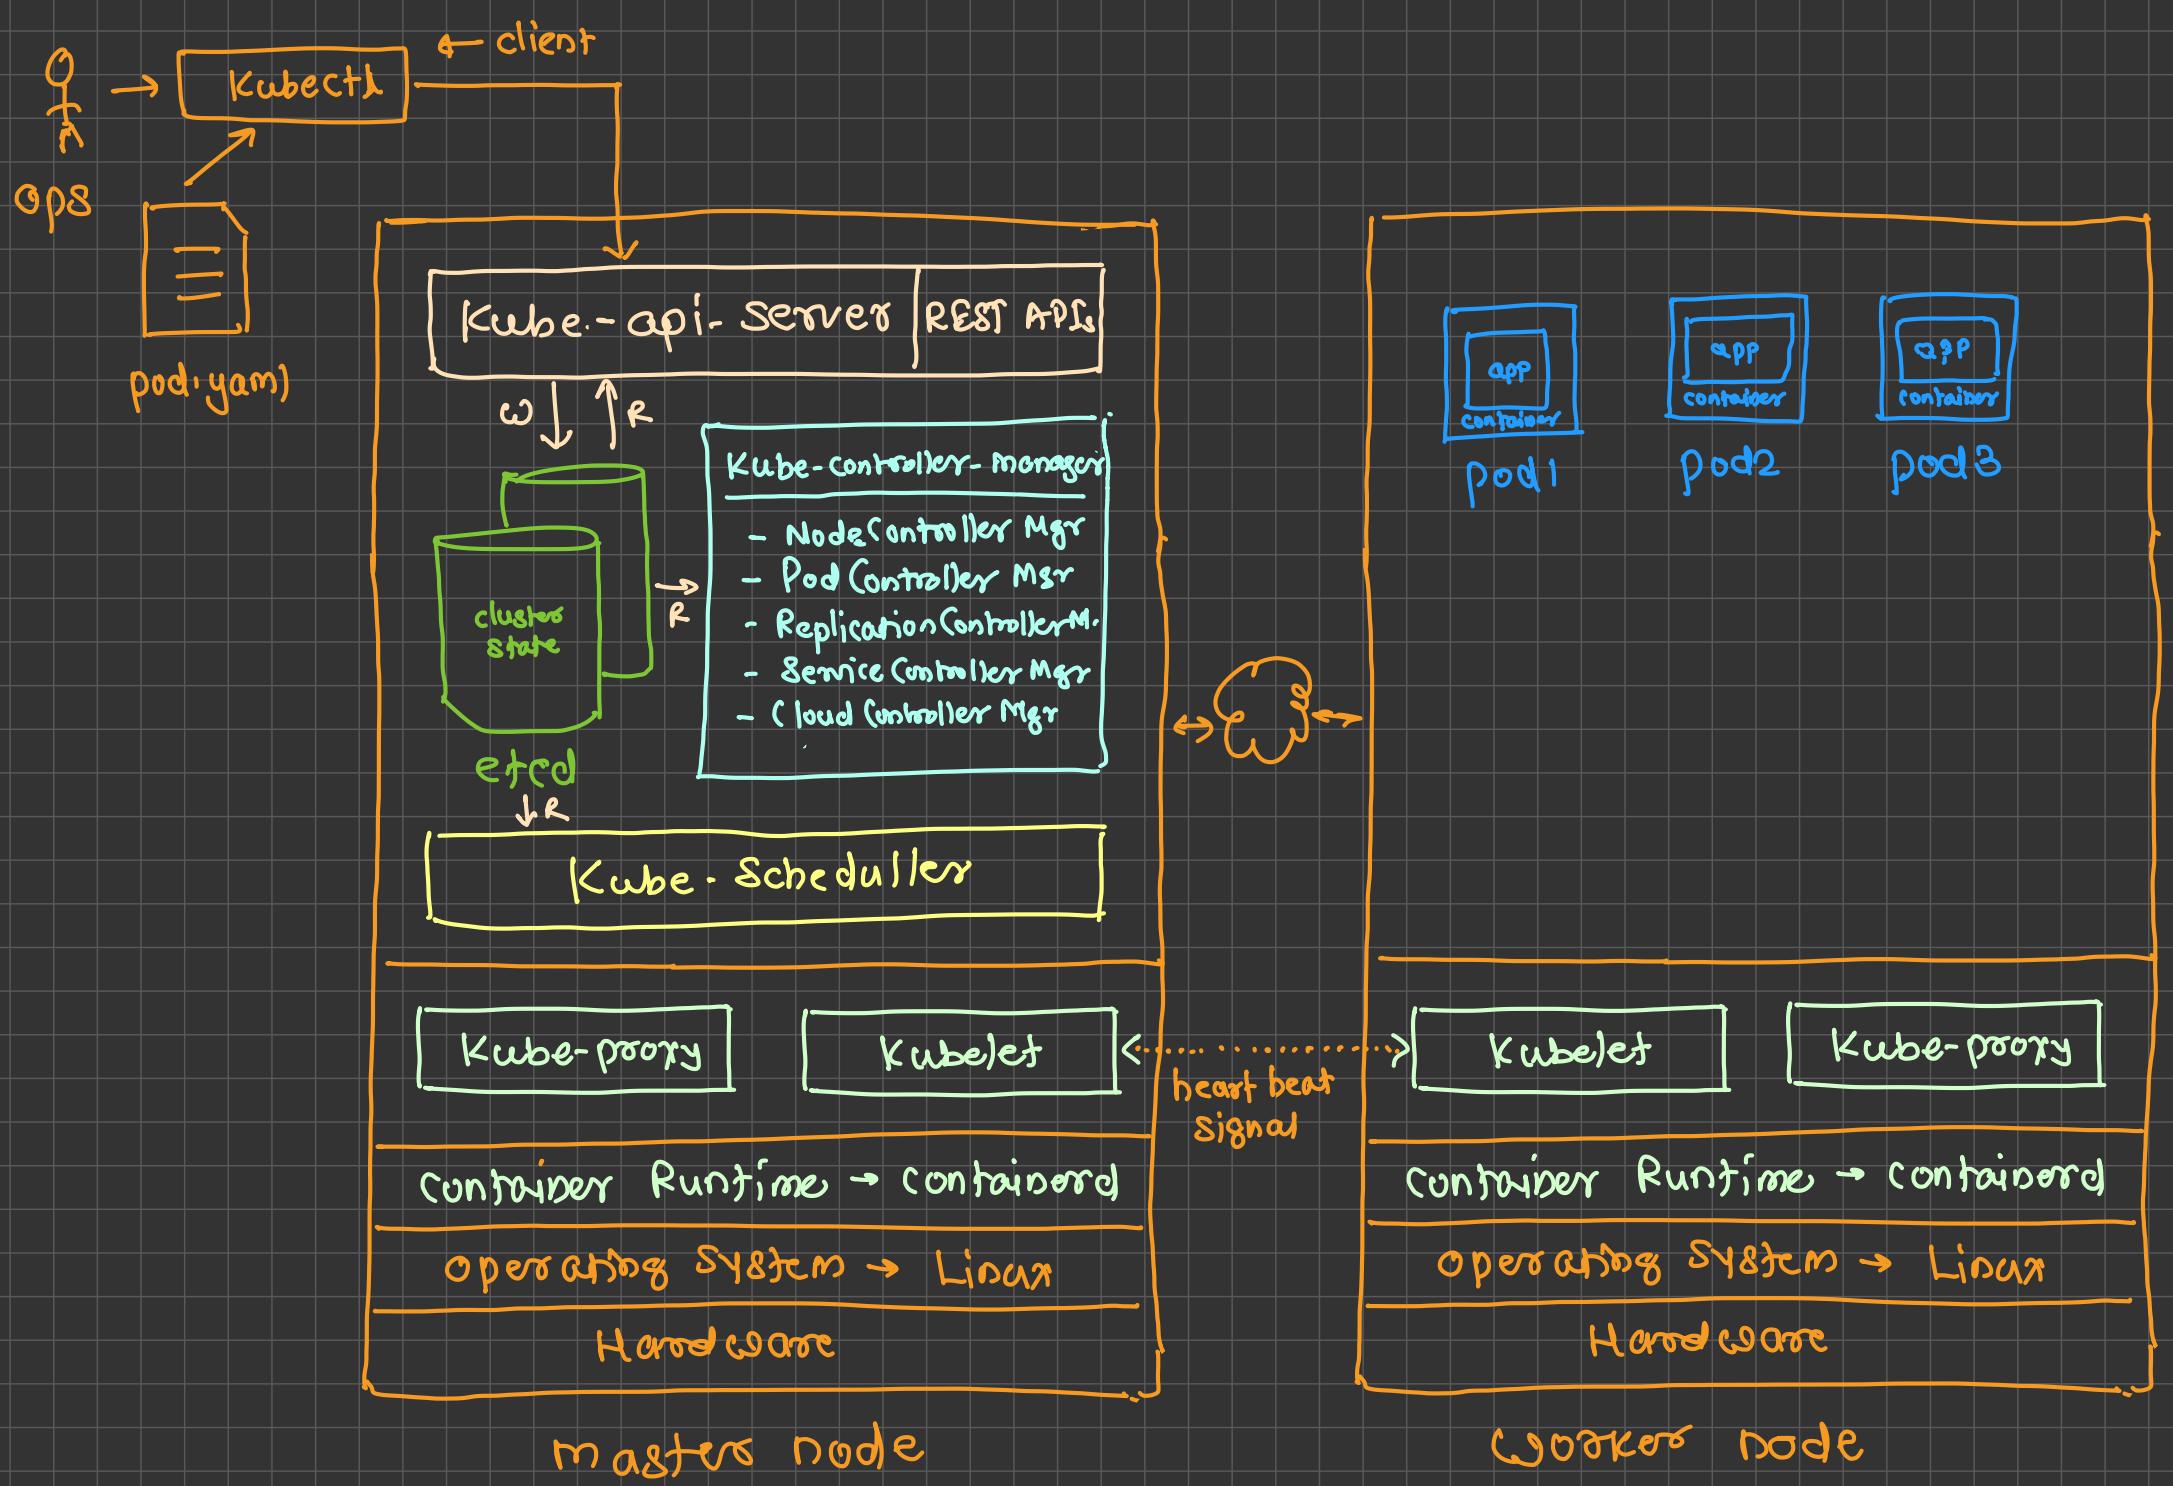
→ multi - master cluster

- more than one masters and more than one workers
- also known HA cluster
- used in pre-prod/ UAT / production env.



Kubernetes Components







Master Components

- These are the **brain** of the Kubernetes cluster — they make all the decisions about what runs where, how scaling happens, and how the system stays healthy.
- Master components can be run on any machine in the cluster

Component	Function
kube-apiserver	The front door of the Kubernetes control plane. It exposes the Kubernetes API , which is used by all components, kubectl, and external clients.
etcd	The database of Kubernetes — stores all cluster data, configurations, and states in a key-value format.
kube-scheduler	Decides which node each Pod should run on based on resource requirements and constraints.
kube-controller-manager	Ensures that the actual state of the cluster matches the desired state defined in manifests.
cloud-controller-manager (optional)	Manages integration between Kubernetes and your cloud provider (e.g., AWS, GCP, Azure).



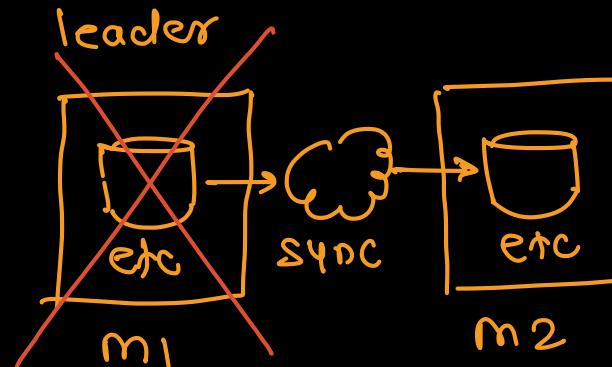
Master Components

■ kube-apiserver — REST APIs

- Acts as the communication hub between users, components, and the cluster.
- Every kubectl command goes through it.
- Validates requests and updates etcd accordingly.
- It's a stateless service – you can run multiple instances for high availability

■ etcd → database

- A distributed, consistent key-value store.
- Stores all cluster data, including:
 - Pod states
 - Configurations
 - Secrets
 - Node information
- It's the source of truth for your cluster.
- It's critical – if etcd is lost, your cluster loses its state.





Master Components

▪ **kube-scheduler**

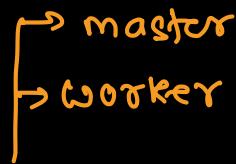
- Watches for new Pods that don't have a Node assigned.
- Chooses the best Node to run the Pod based on:
 - Resource requests (CPU, memory)
 - Node affinity/anti-affinity
 - Taints and tolerations
 - Pod priorities
 - Custom policies

▪ **kube-controller-manager**

- Runs a set of controller loops, each responsible for maintaining part of the system's desired state.
- Examples of controllers:
 - Node Controller – manages node status.
 - Replication Controller – ensures the desired number of pod replicas are running.
 - Endpoint Controller – manages endpoint objects.
 - Service Account & Token Controllers – create default accounts and tokens.



Node Components



- Node components run on every node, maintaining running pods and providing the Kubernetes runtime environment
- **kubelet**
 - Primary agent that runs on every node.
 - Communicates with the API Server.
 - Ensures the containers defined in PodSpecs are running and healthy.
 - Reports node and pod status back to the control plane. → master
 - Watches for Pod definitions assigned to the node and runs them using the container runtime.
- **kube-proxy**
 - Maintains network rules on nodes.
 - Ensures that networking is consistent across the cluster.
 - Implements Kubernetes Service abstraction – enabling communication between different Pods and Services.
 - Uses iptables or IPVS to forward traffic to the correct Pod endpoints.
 - Supports load balancing between Pods behind a Service.
- **Container Runtime**
 - The actual software responsible for running containers.
 - The kubelet interacts with the runtime through the Container Runtime Interface (CRI).
 - Common runtimes:
 - containerd (default on most modern clusters)
 - CRI-O
 - Docker Engine (deprecated as of K8s v1.24+)
 - Mirantis Container Runtime



Create Cluster

- Use following commands on both master and worker nodes

```
> sudo apt-get update && sudo apt-get install -y apt-transport-https curl  
> curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -  
> cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list deb https://apt.kubernetes.io/kubernetes-xenial main EOF  
> sudo apt-get update  
> sudo apt-get install -y kubelet kubeadm kubectl  
> sudo apt-mark hold kubelet kubeadm kubectl
```



Initialize Cluster Master Node

- Execute following commands on master node

```
> kubeadm init --apiserver-advertise-address=<ip-address> --pod-network-cidr=10.244.0.0/16  
> mkdir -p $HOME/.kube  
> sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
> sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Install pod network add-on

```
> kubectl apply -f  
https://raw.githubusercontent.com/coreos/flannel/2140ac876ef134e0ed5af15c65e414cf26827915/Documentation/kube-flannel.yml
```



Add worker nodes

- Execute following command on every worker node

```
> kubeadm join --token <token> <control-plane-host>:<control-plane-port> --discovery-token-ca-cert-hash sha256:<hash>
```



Steps to install Kubernetes

5

initialize

4

POD Network

3

kubeadm

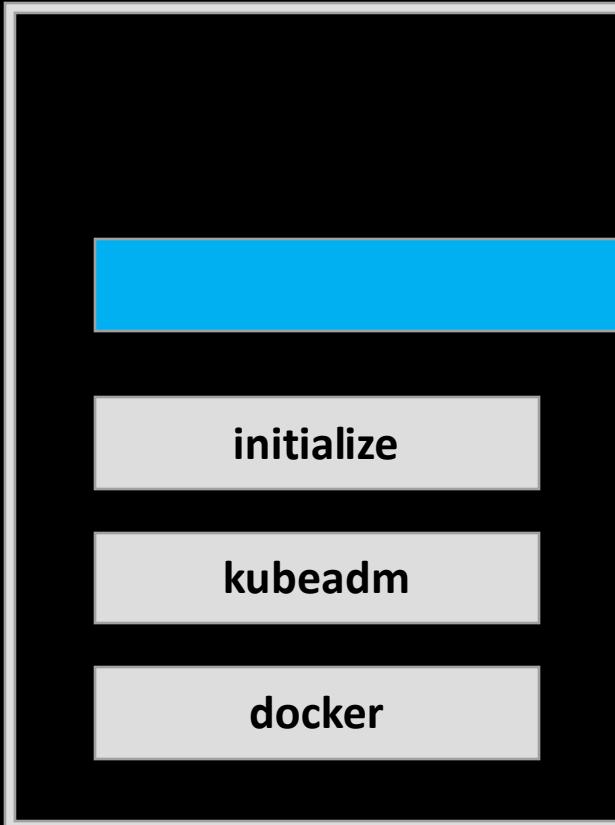
2

docker

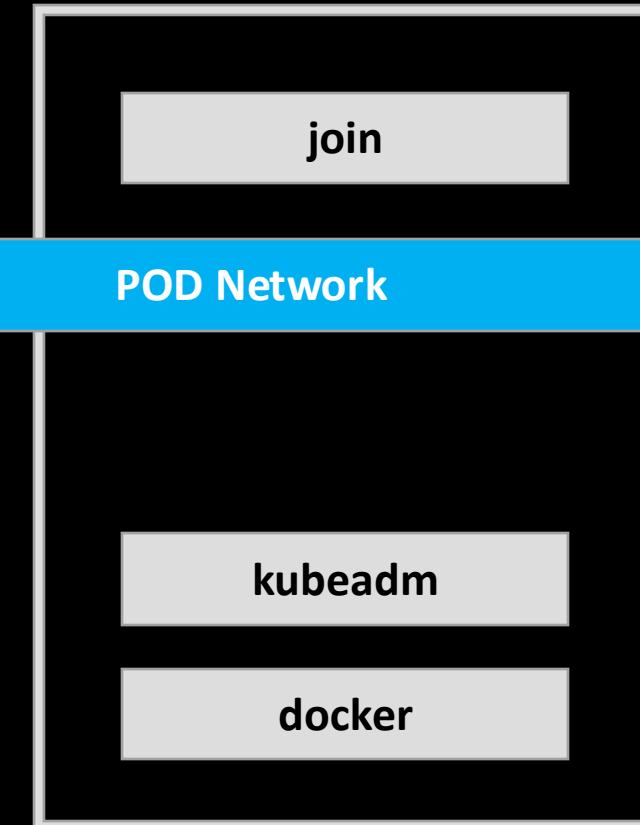
1

join

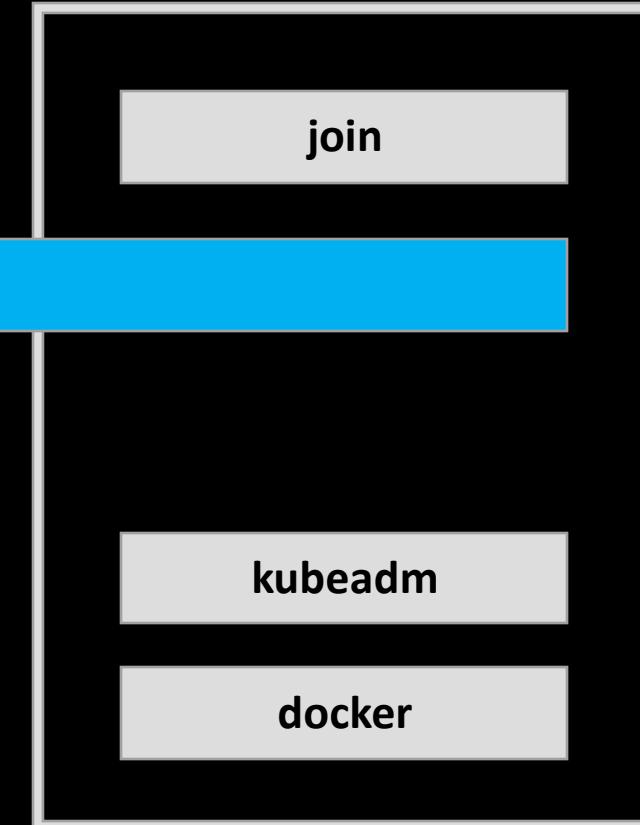
join



Master



Worker Node 1



Worker Node 2



Kubernetes Objects

- The basic Kubernetes objects include

- Pod
- Service
- Volume
- Namespace

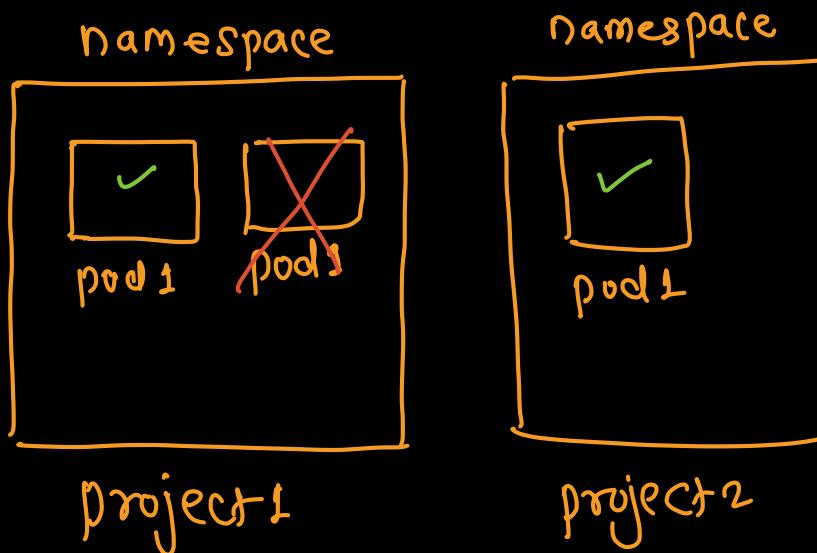
- Kubernetes also contains higher-level abstractions build upon the basic objects

- Deployment
- DaemonSet
- StatefulSet
- ReplicaSet
- Job



Namespace

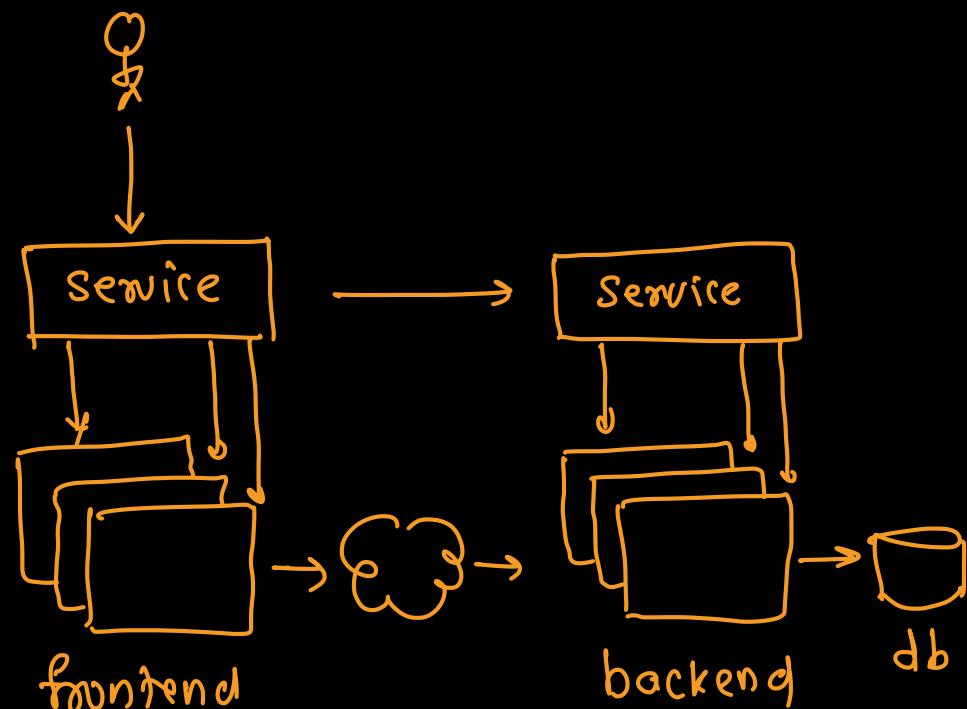
- Namespaces are intended for use in environments with many users spread across multiple teams, or projects
- Namespaces provide a scope for names
- Names of resources need to be unique within a namespace, but not across namespaces
- Namespaces can not be nested inside one another and each Kubernetes resource can only be in one namespace
- Namespaces are a way to divide cluster resources between multiple users / projects / team

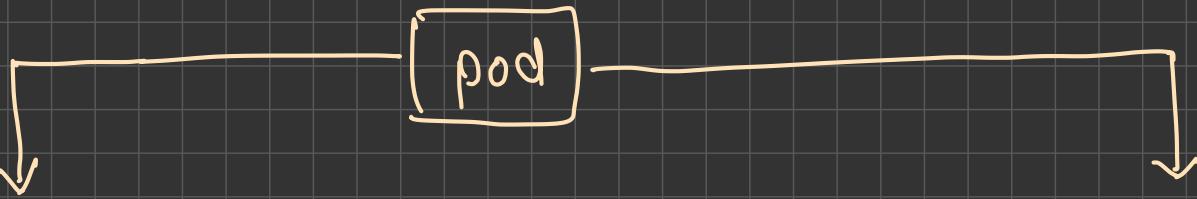




Pod

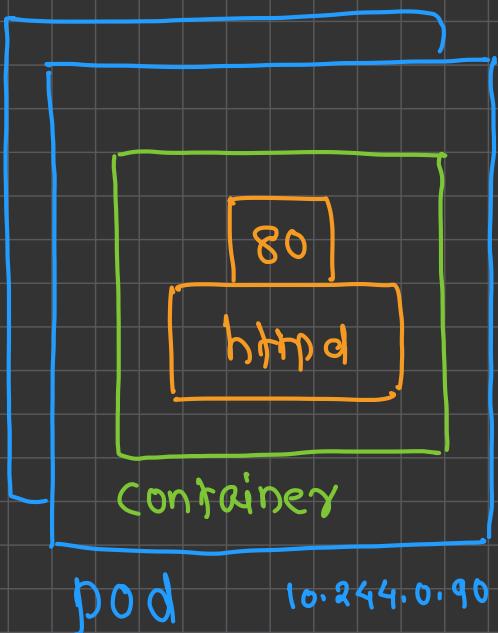
- A Pod is the basic execution unit of a Kubernetes application
- The smallest and simplest unit in the Kubernetes object model that you create or deploy
- A Pod represents processes running on your Cluster
- Pod represents a unit of deployment
- A Pod encapsulates
 - application's container (or, in some cases, multiple containers)
 - storage resources
 - a unique network IP
 - options that govern how the container(s) should run
 - ↳ configuration → resources → CPU/memory





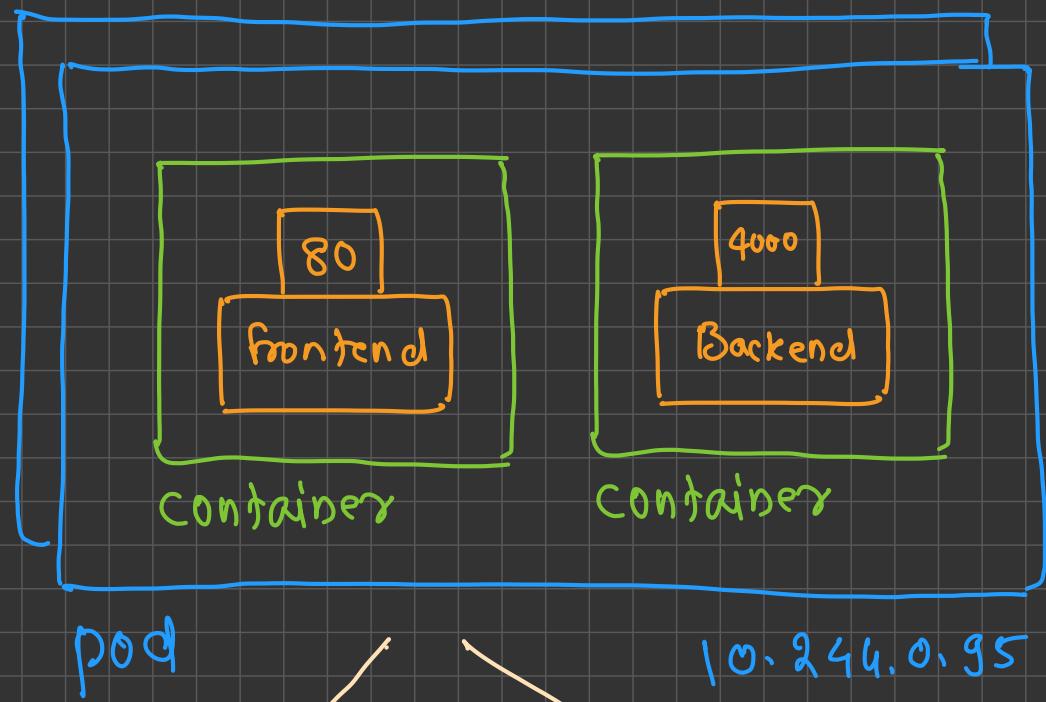
Single container pod

multi-containers pod



Container

10.244.0.90



Container

Container

10.244.0.95

Sidecar pattern

init container pattern

YAML → YAML An't markup Language → recursive acronymes

20, person), Karad, p@test.com ← raw data → meaningless data



serialization format



meaningful information

XML

```
<person>
  <age>20</age>
  <name>person</name>
  <city>Karad</city>
  ;
</person>
```

JSON

```
{ "age": 20,
  "name": "person",
  "city": "Karad",
  :
```

YAML

```
age: 20
name: "person"
city: 'Karad'
email: p1@test.com
```

YAML

- case sensitive → city & City are different
- tabs are NOT allowed
- Spaces are used instead of tabs
- use .yaml or .yml as file extension
- space is required after : & -
- data types

→ scalar

- single values → 60, person1, karad
- " " are NOT required for string values

→ map

→ pair of key-value

→ e.g. name: person1, age: 20

→ list

→ collection of values (scalar or map)

→ e.g. # languages

- C

- C++

- python

| - name: person1
| age: 20
| - name: person2
| age: 30



YAML to create Pod

apiVersion: v1

kind: Pod

metadata:

name: myapp-pod

labels:

app: myapp

spec:

containers:

- name: myapp-container

image: httpd

apiVersion

→ version of k8s api

→ for basic objects like pod, service etc = v1

→ for higher level objects like Replicaset, Deployment = apps/v1

kind

→ type of object/resource to be created

→ e.g. Pod, Service, Deployment etc.

metadata

→ more information of object to be created

→ e.g. name, namespace, labels

Spec

→ Specification of object to be created

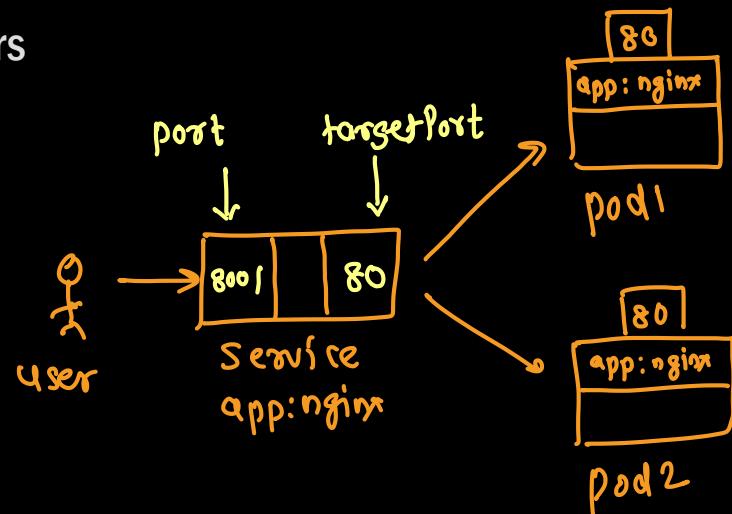
→ depends on the object

Service → Load balancer



- An abstract way to expose an application running on a set of Pods as a network service
- Service is an abstraction which defines a logical set of Pods and a policy by which to access them (sometimes this pattern is called a micro-service)
- Service Types
 - ClusterIP
 - LoadBalancer
 - NodePort

↳ load balancing

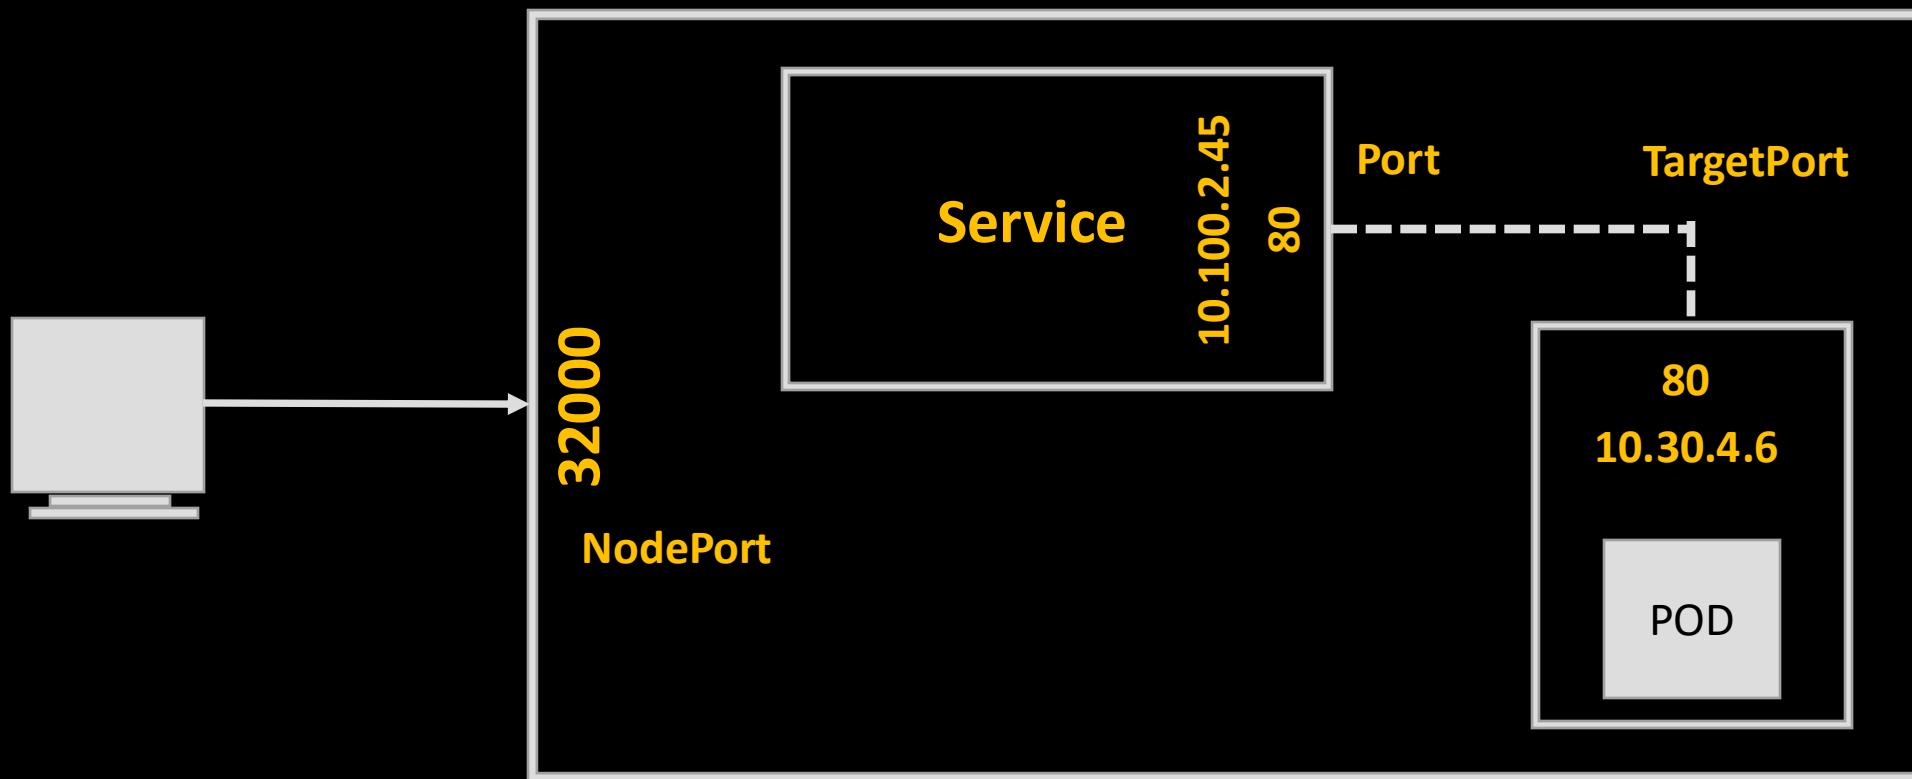


```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector: {label configuration on pod(s)}
  app: MyApp
  ports:
    - protocol: TCP
      port: 80 ← User sends request
      targetPort: 8080 80
      ↳ pod's port
```



Service Type: NodePort

- Exposes the Service on each Node's IP at a static port (the NodePort)
- You'll be able to contact the NodePort Service, from outside the cluster, by requesting <NodeIP>:<NodePort>





Replica Set

label → map (Key-value pair)

desired count

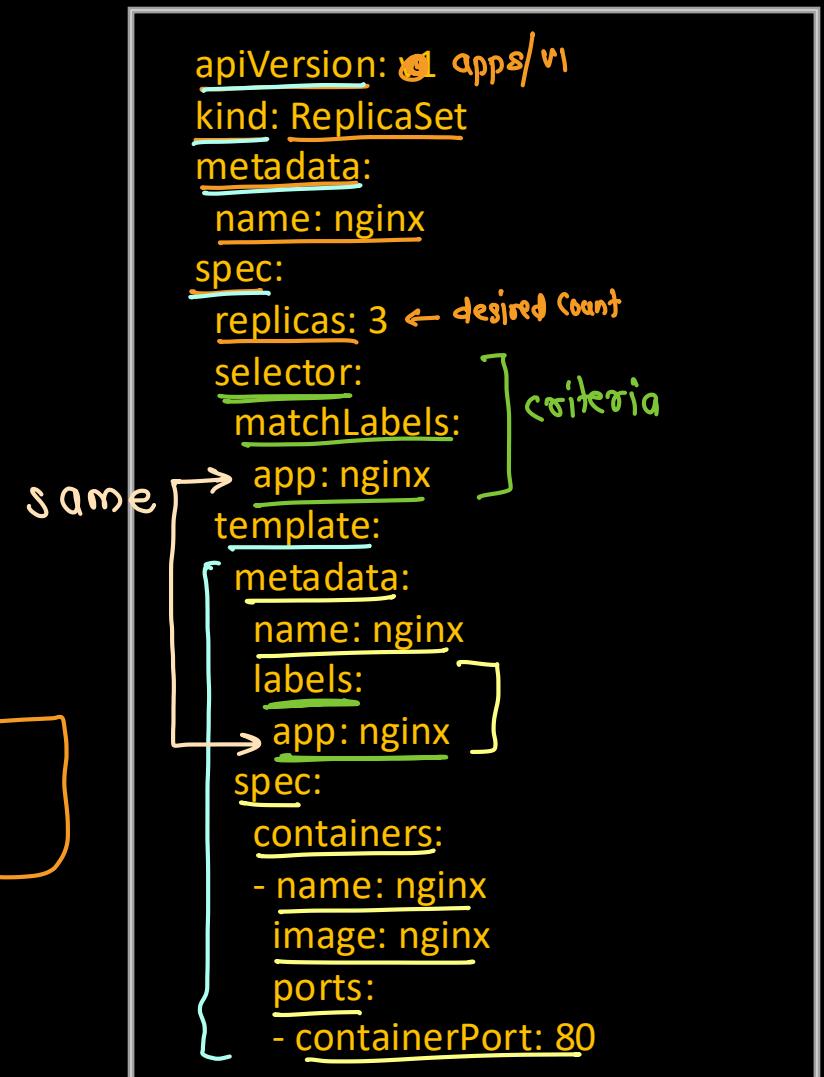
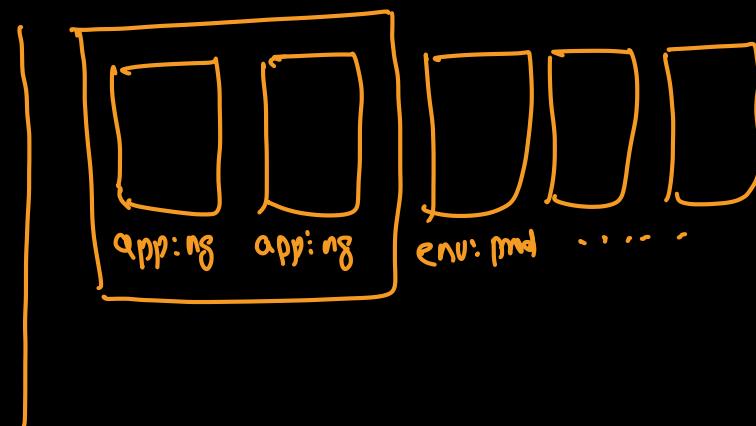
- A Replica Set ensures that a specified number of pod replicas are running at any one time
- In other words, a Replica Set makes sure that a pod or a homogeneous set of pods is always up and available
- If there are too many pods, the Replica Set terminates the extra pods
- If there are too few, the Replica Set starts more pods
- Unlike manually created pods, the pods maintained by a Replica Set are automatically replaced if they fail, are deleted, or are terminated

Docker swarm →

Container Scaling → Service
load balancing → service

Kubernetes →

Container Scaling → Replicaset
load balancing → service





Deployment

- A Deployment provides declarative updates for Pods and ReplicaSets
- You describe a *desired state* in a Deployment, and the Deployment Controller changes the actual state to the desired state at a controlled rate
- You can use deployment for
 - Rolling out ReplicaSet
 - Declaring new state of Pods
 - Rolling back to earlier deployment version
 - Scaling up deployment policies
 - Cleaning up existing ReplicaSet

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: website-deployment
spec:
  selector:
    matchLabels:
      app: website
  replicas: 10
  template:
    metadata:
      name: website-pod
      labels:
        app: website
    spec:
      containers:
        - name: website-container
          image: pythoncpp/test_website
      ports:
        - containerPort: 80
```